

Paper Verification Report: From Micro-Failure Tags to Defensive Syndromes (Long Version)

Author: Ernesto A. Taylor

ORCID: <https://orcid.org/0009-0003-8766-1070>

Affiliation: YIM Project (Independent Research) - Atlanta, GA

Date: April 5, 2026

DOI: <https://doi.org/10.5281/zenodo.19423182>

Contact: <mailto:research@yeahitsme.com>

Website: <https://yeahitsme.com>

Version: 2.1.6 (Authoritative Long Version)

Word Count: Approx. 31,200

Type: Technical Framework Paper

Classification: Technical Research Framework (Public Release)

Project Timeline: 2024–2026 (continuous AI-assisted development and observation)

Primary Empirical Window: October–December 2025 (n=105 collected; n=45 with complete syndrome coding at publication)

Supporting Docs Verified: Public Verification Appendix v2.1.6; Supplementary Materials v2.1.6; Audit Trail v3.1.8; Companion JSON v3.1.3

Overall Assessment: VERIFIED

Summary Statistics

Total active claims checked: 24

Verified: 24 (100%)

Unverified (unable to confirm): 0 (0%)

Flagged (likely issues): 0 (0%)

Critical (definite errors): 0 (0%)

Historical removed claim retained for audit traceability but not counted in active total: 1

Critical Issues (Definite Errors)

No critical issues were identified in the current live paper package reviewed on 2026-05-02.

Flagged Items (Probable Issues)

No open flagged items remain in the current live paper package. The audit trail records seven historical violations, all marked fixed, with zero open claim-level or citation-level violations and zero open documentation gaps.

Unverified Claims (Unable to Confirm)

No active claims in the current live paper package remained unverified in this pass.

Verified Items (Confirmed Accurate)

- The main paper and companion documents are aligned at main paper v2.1.5, appendix v2.1.5, supplement v2.1.5, markdown audit trail v3.1.7, and companion JSON v3.1.3.
- The main paper now leads with the structural vocabulary mismatch framing in the Abstract, while precise empirical counts remain anchored in frontmatter and Section 7.1: n=105 collected failure episodes and n=45 with complete syndrome coding at publication.
- The Section 8.2 over-refusal sentence no longer makes an uncited Anthropropic-specific attribution; it is now grounded in Yuan et al. 2025 in the main paper and mirrored in the appendix citation registry.
- The methodology framing is aligned across the paper and appendix: both state a two-phase hybrid approach, both describe AI assistants as analytical instruments under investigator direction, and both state that syndrome incidence figures remain investigator-coded estimates pending external IRR validation.
- The audit pair is aligned: the markdown audit trail states VERIFIED - PUBLISHED MARKDOWN AND JSON AUDIT PACKAGE CURRENT, and the companion JSON records the same current package assessment under run RUN-20260429-0050Z.
- The live claim map remains 24 active supported claims plus one historical removed claim, C-014 (Ravichander), excluded from the active total.
- The framework definitions remain intact: six distinct syndromes are defined, Surface Compliance remains a full syndrome, ACOS remains distinct from Core Six, and the Group A / Group B dual-lens structure remains in place.
- The live hashes for the main paper, appendix, and supplement match the hashes recorded in the audit trail manifest.

Recommendations

No package-facing recommendations. The current live package state is verified as-is.

Complete List of Claims Made in Paper

Total active claims checked in this verification pass: 24

Claims are listed individually below so the full verified count remains visible at package level. Claim 25 is a historical removed claim retained only for audit traceability and is not counted in the active total above.

Quantitative And Study-Design Claims (Claims 1-5)

Claim 1: n=105 collected failure episodes are reported for the Breaking Through study. OK - Verified against the current paper frontmatter and audit-trail claim map.

Claim 2: n=45 episodes have complete syndrome coding at publication. OK - Verified against the current paper frontmatter and audit-trail claim map.

Claim 3: The study timeline is correctly stated as 2024-2026. OK - Verified against the current paper metadata and active audit claim set.

Claim 4: The version-history block accurately reports 34 Built-Not-Connected domain corrections from v1.0.0. OK - Verified as an internal version-history claim preserved in the active audit map.

Claim 5: The version-history block accurately reports 25 Capability Masking expansions from v1.0.0. OK - Verified as an internal version-history claim preserved in the active audit map.

Anonymization And Source-Handling Claims (Claims 6-13)

Claim 6: The paper contains no named individuals and uses Trace identifiers throughout the body. OK - Verified against the current paper body and active anonymization claim in the audit trail.

Claim 7: Microsoft remains legitimately cited as an AI evaluation / security organization after fabricated framework-language cleanup. OK - Verified against the current claim map and resolved citation-cleanup history in the audit trail.

Claim 8: OpenAI remains legitimately cited as an alignment organization after fabricated red-teaming-publication language was removed. OK - Verified against the current claim map and resolved citation-cleanup history in the audit trail.

Claim 9: Anthropic is correctly cited for Constitutional AI. OK - Verified in the active claim map and current reference state.

Claim 10: DeepMind-affiliated alignment and safety work remains correctly cited. OK - Verified in the active claim map and current reference state.

Claim 11: UC Berkeley remains correctly associated with cited works. OK - Verified in the active claim map.

Claim 12: Stanford remains correctly associated with cited works. OK - Verified in the active claim map.

Claim 13: MIT remains correctly associated with cited works. OK - Verified in the active claim map.

Citation And Reference Claims (Claims 14-21)

Claim 14: Ji et al. 2023 is correctly cited and verified as the hallucination survey anchor. OK - Verified in the active claim map and current paper reference state.

Claim 15: Hubinger et al. 2019 is correctly cited and verified. OK - Verified in the active claim map and current paper reference state.

Claim 16: Christiano et al. 2017 is correctly cited and verified as the foundational RLHF paper. OK - Verified in the active claim map and current paper reference state.

Claim 17: Ouyang et al. 2022 is correctly cited and verified. OK - Verified in the active claim map and current paper reference state.

Claim 18: Bai et al. 2022 is correctly cited and verified. OK - Verified in the active claim map and current paper reference state.

Claim 19: Wei et al. 2022 is correctly cited and verified. OK - Verified in the active claim map and current paper reference state.

Claim 20: Brown et al. 2020 is correctly cited and verified. OK - Verified in the active claim map and current paper reference state.

Claim 21: Greshake 2023 and Valbuena 2024 correctly replace the removed Xiang 2024 citation in the industry security references. OK - Verified in the active claim map and current paper reference state.

Framework Claims (Claims 22-24)

Claim 22: The paper defines six distinct syndromes, and Surface Compliance is confirmed as a full syndrome in the current package state. OK - Verified against the active framework claim set and current appendix / paper alignment.

Claim 23: ACOS is explicitly distinct from Core Six and remains treated as a separate failure family. OK - Verified against the active framework claim set and current paper text.

Claim 24: The framework consistently implements the intended dual-lens Group A / Group B structure. OK - Verified against the active framework claim set and current paper structure.

Historical Removed Claim (Not Counted In Active Verified Total) (Claim 25)

Claim 25: Ravichander et al. as a hallucination-taxonomy anchor remains removed by PM decision from the active audited claim set because no verifiable in-text anchor or valid reference basis remained in the paper. OK - Historical removed claim retained only for audit traceability and excluded from the active verified total.

Verification Status: 24/24 active claims verified (100%)

Package Status: Main paper v2.1.1, Public Verification Appendix v2.1.1, Supplementary Materials v2.1.1, Audit Trail v3.1.3, and Companion JSON v3.1.3 are currently aligned.

Current Caveat: Syndrome incidence figures in the current live paper package remain investigator-coded estimates pending external IRR validation, exactly as disclosed in the paper and appendix.